

## Solutions for Exercise 1

### Import libraries

```
In [2]: import numpy as np
import matplotlib.pyplot as plt
import matplotlib inline

import sklearn
from sklearn.datasets import load_digits
from sklearn import model_selection
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
```

## 1 Exploring the Data

```
In [3]: # Load digits data set
digits = load_digits()

data = digits["data"]
images = digits["images"]
target = digits["target"]
target_names = digits["target_names"]

print(f"data.shape = {data.shape}")
print(f"data.dtype = {data.dtype}")
print(f"images.shape = {images.shape}")
print(f"images.dtype = {images.dtype}")
print(f"target.shape = {target.shape}")
print(f"target.dtype = {target.dtype}")
print(f"target_names.shape = {target_names.shape}")
print(f"target_names.dtype = {target_names.dtype}")
print(f"target[:20] = {target[:20]}")

img = images[3, :, :]
assert 2 == len(img.shape)

fig, axes = plt.subplots(1, 3, figsize=(10, 3), tight_layout=True)

axes[0].imshow(img, cmap="gray", interpolation="nearest")
axes[0].set_axis_off()
axes[1].imshow(img, cmap="gray", interpolation="bilinear")
axes[1].set_axis_off()
axes[2].imshow(img, cmap="gray", interpolation="bicubic")
axes[2].set_axis_off()

plt.show()

data.shape = (1797, 64)
data.dtype = float64
images.shape = (1797, 8, 8)
images.dtype = float64
target.shape = (1797,)
target.dtype = int64
target_names.shape = (10,)
target_names.dtype = str
target_names[:20] = [0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9]
```

The right-most image illustrates why we generally use `interpolation = 'nearest'`. For scientific analysis, it's important that we can see the exact value of each pixel in an image array, even if interpolation = 'bicubic' often yields visually more pleasing results.

If your image has the shape `(width, height)` `matplotlib` will display it using a colormap, `viridis` by default. You can change the `cmap` parameter to get any colormap you want, this parameter is ignored for RGB data (images of shape `(width, height, 3)`).

```
In [4]: """
This function filters the digits (3, 9) from the dataset and randomly splits it in train and test set.
"""
# Load data
digits = load_digits()

data = digits["data"]
target = digits["target"]

# Data filtering
num_1, num_2 = 3, 9
mask = np.logical_or(target == num_1, target == num_2)
data = data[mask]
target = target[mask]

# Relabel targets
target[target == num_1] = -1
target[target == num_2] = 1

# Split into train and test data
X_all = data
y_all = target
X_train, X_test, y_train, y_test = model_selection.train_test_split(
    X_all, y_all, test_size=0.4, random_state=0)
print(f"X_train.shape = {X_train.shape}")
print(f"X_test.shape = {X_test.shape}")
print(f"y_train.shape = {y_train.shape}")
print(f"y_test.shape = {y_test.shape}")
print(f"y_test[:10] = {y_test[:10]}")

X_train.shape = (217, 64)
X_test.shape = (146, 64)
y_train.shape = (217,)
y_test.shape = (146,)
y_test[:10] = [ 1  1  1 -1 -1 -1 -1 -1 -1 -1]
```

## 2 Hand-crafted classifier

### 2.1 Feature construction

Visualize some images to get a sense of which features might be important to distinguish the digits.

```
In [5]: fig, axes = plt.subplots(1, 4, figsize=(12, 3), tight_layout=True)

index = np.random.randint(0, 100)
for i in range(4):
    axes[i].imshow(X_train[y_train==1][index+i].reshape(8,8), cmap="gray", interpolation="nearest")
fig, axes = plt.subplots(1, 4, figsize=(12, 3), tight_layout=True)

index = np.random.randint(0, 100)
for i in range(4):
    axes[i].imshow(X_train[y_train==1][index+i].reshape(8,8), cmap="gray", interpolation="nearest")

0 0 0 0 0
1 1 1 1 1
2 2 2 2 2
3 3 3 3 3
4 4 4 4 4
5 5 5 5 5
6 6 6 6 6
7 7 7 7 7
0 2 4 6 8
```

```
In [6]: def features2d(x):
"""
Perform a user defined dimension reduction
Input: #instances x 64 numpy array
Output: #instances x 2 numpy array
"""
images = x.reshape(-1, 8, 8) #instances x row index x column index
mu = np.mean(images[:, 2:4, 2:4], axis=(1, 2))
feat2 = np.mean(images[:, 4:6, 4:7], axis=(1, 2))
return np.array([feat1, feat2]).T

def worse_features2d(x):
"""
Perform a user defined dimension reduction
Input x: #instances x 64 numpy array
Output: #instances x 2 numpy array
"""
# mean(image)
feat1 = np.mean(x, axis=-1)

# var(image)
feat2 = np.var(x, axis=-1)

return np.array([feat1, feat2]).T
```

The `features2d()` function looks at the average pixel values of two patches in the image. They are informative enough to more or less differentiate between the digit '3' and '9' which we can see in the scatter plot below.

The `worse_features2d()` is much simpler and doesn't extract good features. We won't be using it for the rest of the exercise.

### 1.2 Scatter Plot

```
In [7]: def scatter_plot_simple(x, y, title="Training"):
"""
This function returns the dataset scatter plot
Input x: N x 2 numpy array
Input y: N x 1 numpy array
Output: None
"""
plt.figure(figsize=(8,8))
plt.gca().set_aspect('equal')
plt.title(title + " Scatter Plot: 3 vs 9")

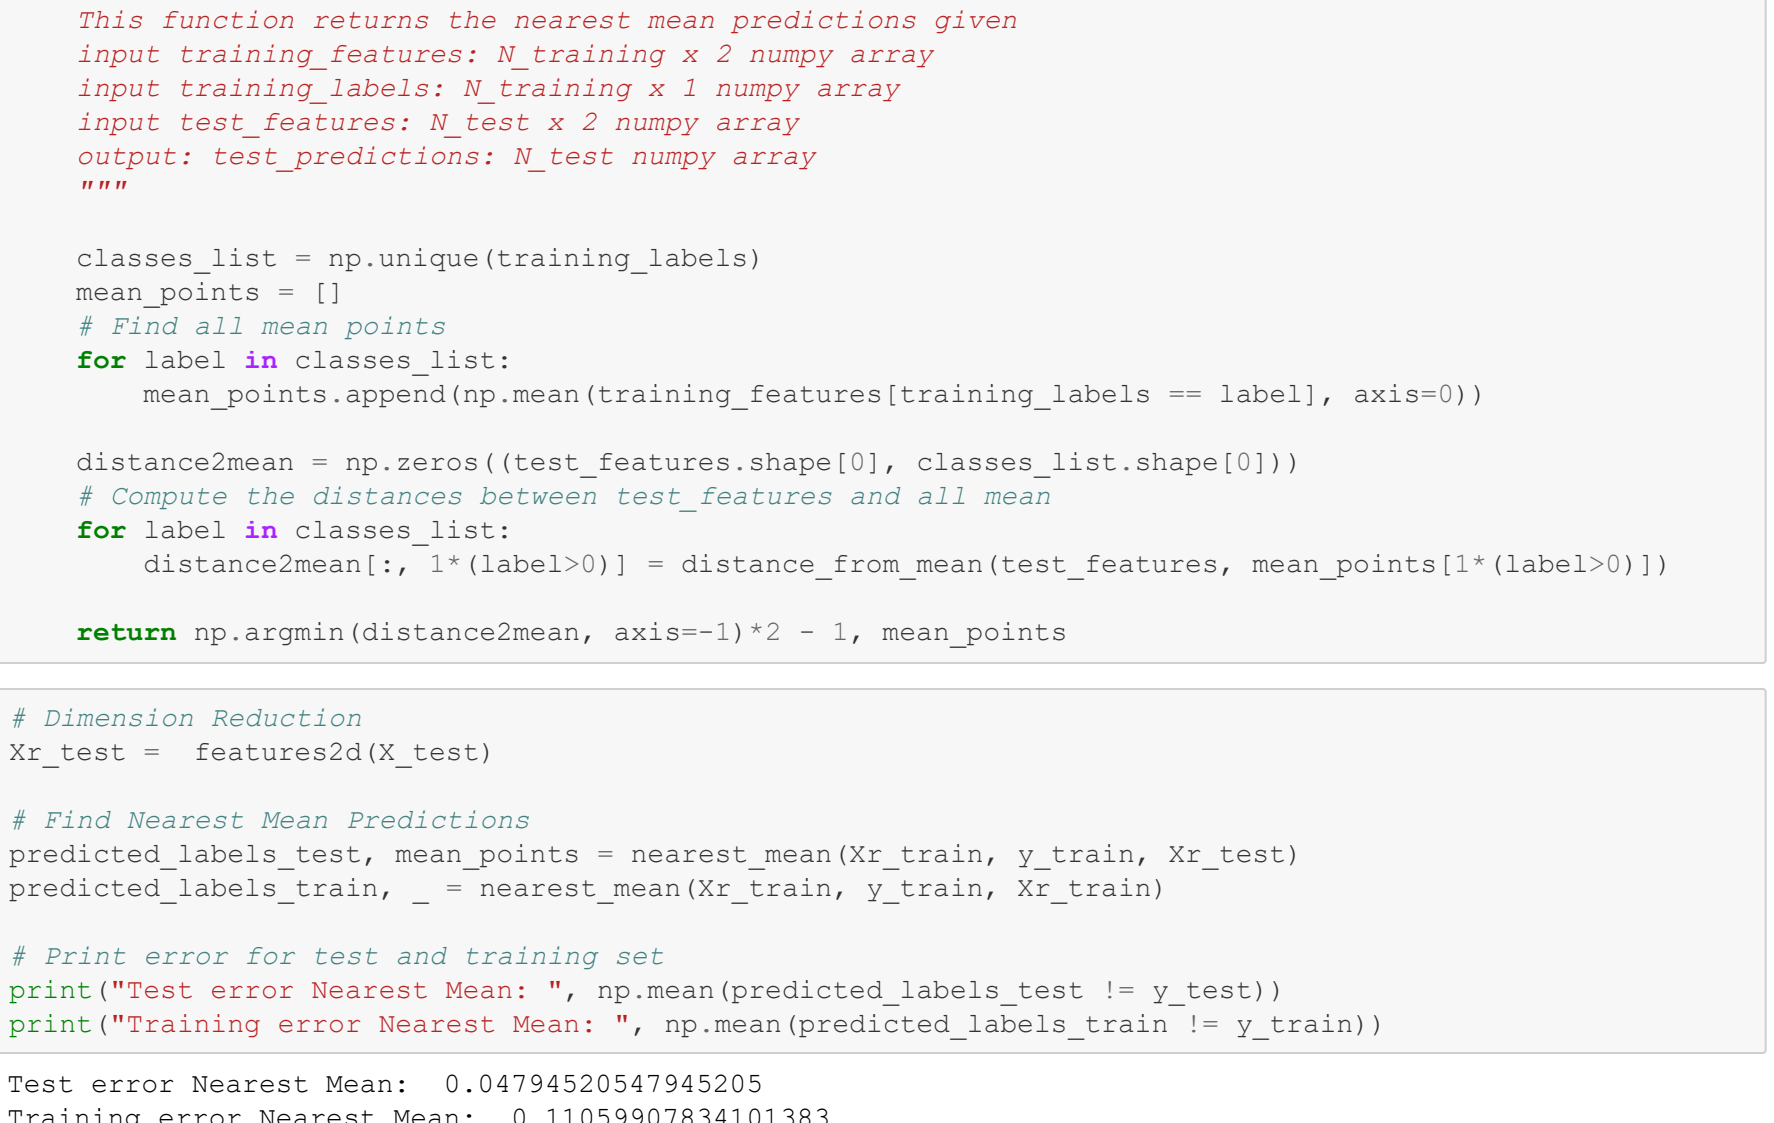
# Scatter plot
plt.scatter(x[:, -1, 0], x[:, -1, 1], marker="o", s=30, c="b", label="3")
plt.scatter(x[:, 1, 0], x[:, 1, 1], marker="x", s=30, c="r", label="9")

plt.xlabel("feature 1")
plt.ylabel("feature 2")
plt.legend()
plt.show()
```

```
In [8]: # Load data
digits = load_digits()

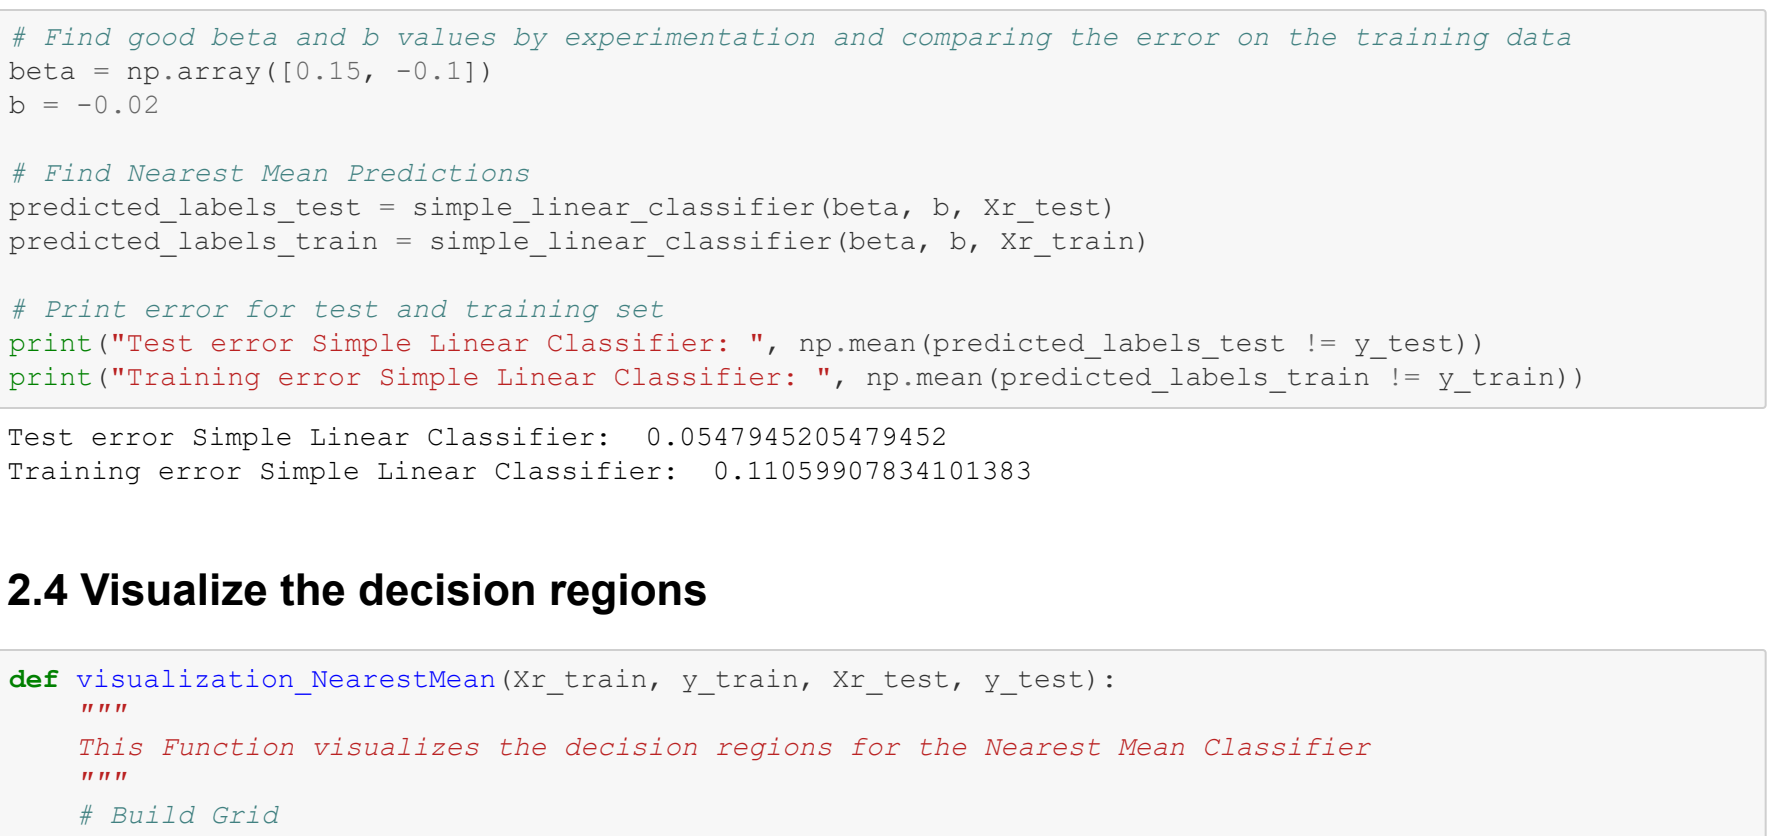
# Dimension Reduction
Xr_train, Xr_test = features2d(X_train), features2d(X_test)

# Scatter Plot
scatter_plot_simple(Xr_train, y_train, "Training")
scatter_plot_simple(Xr_test, y_test, "Test")
```



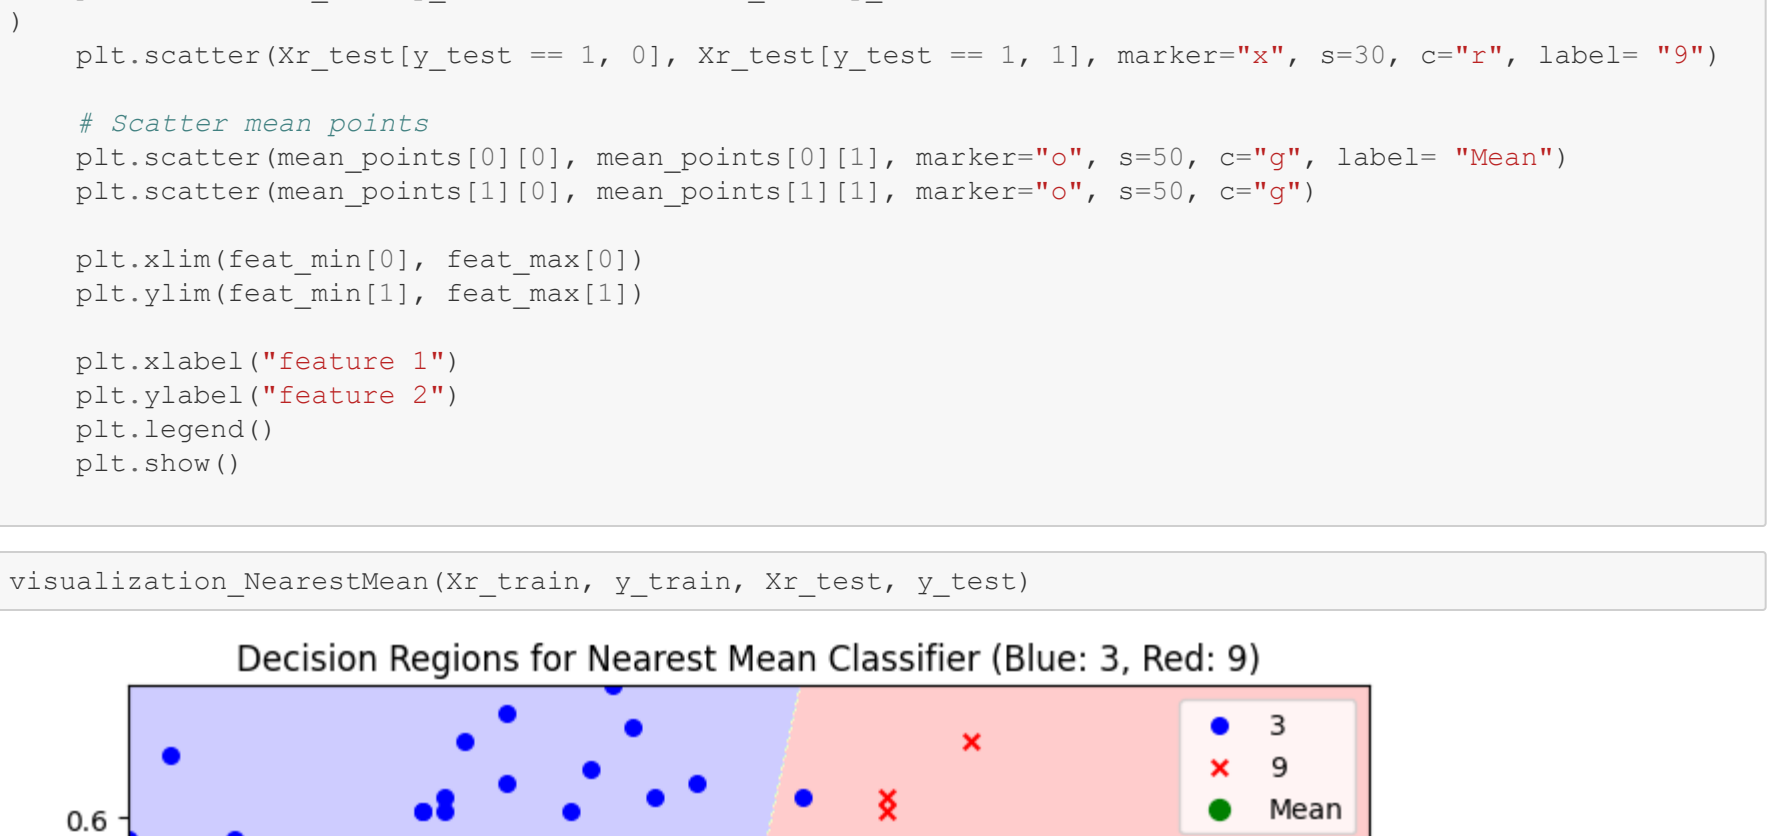
```
In [9]: # Dimension Reduction
_Xr_train, _Xr_test = worse_features2d(X_train), worse_features2d(X_test)

# Scatter Plot
scatter_plot_simple(_Xr_train, y_train, "Training - Worse features")
```



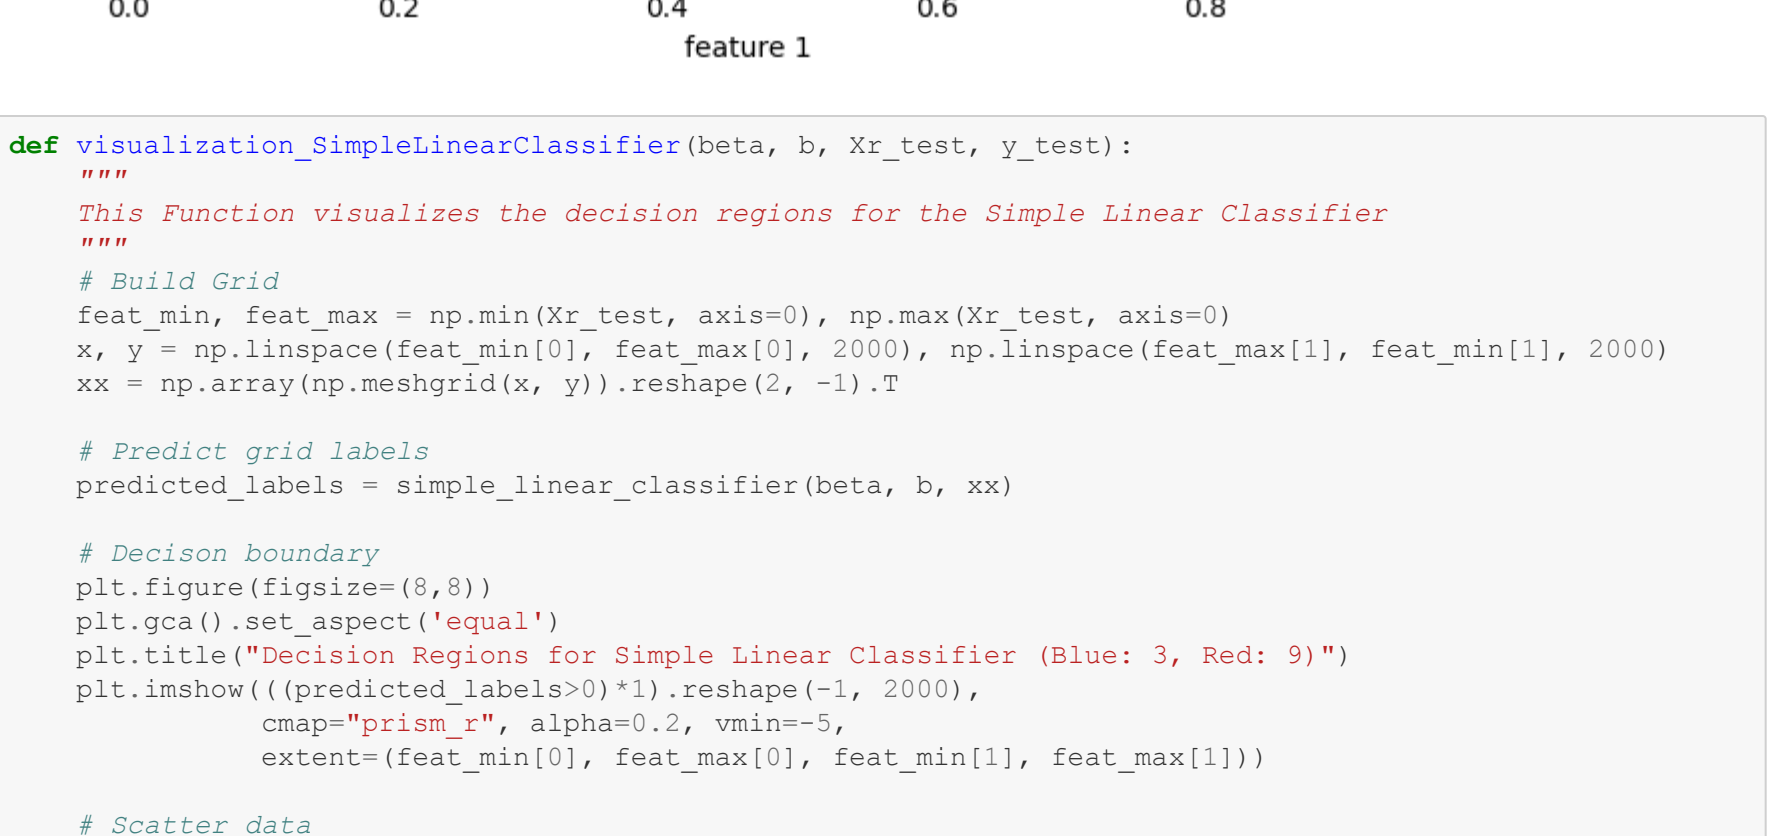
```
In [9]: # Dimension Reduction
_Xr_train, _Xr_test = worse_features2d(X_train), worse_features2d(X_test)

# Scatter Plot
scatter_plot_simple(_Xr_train, y_train, "Training - Worse features")
```




```
In [9]: # Dimension Reduction
_Xr_train, _Xr_test = worse_features2d(X_train), worse_features2d(X_test)

# Scatter Plot
scatter_plot_simple(_Xr_train, y_train, "Training - Worse features")
```



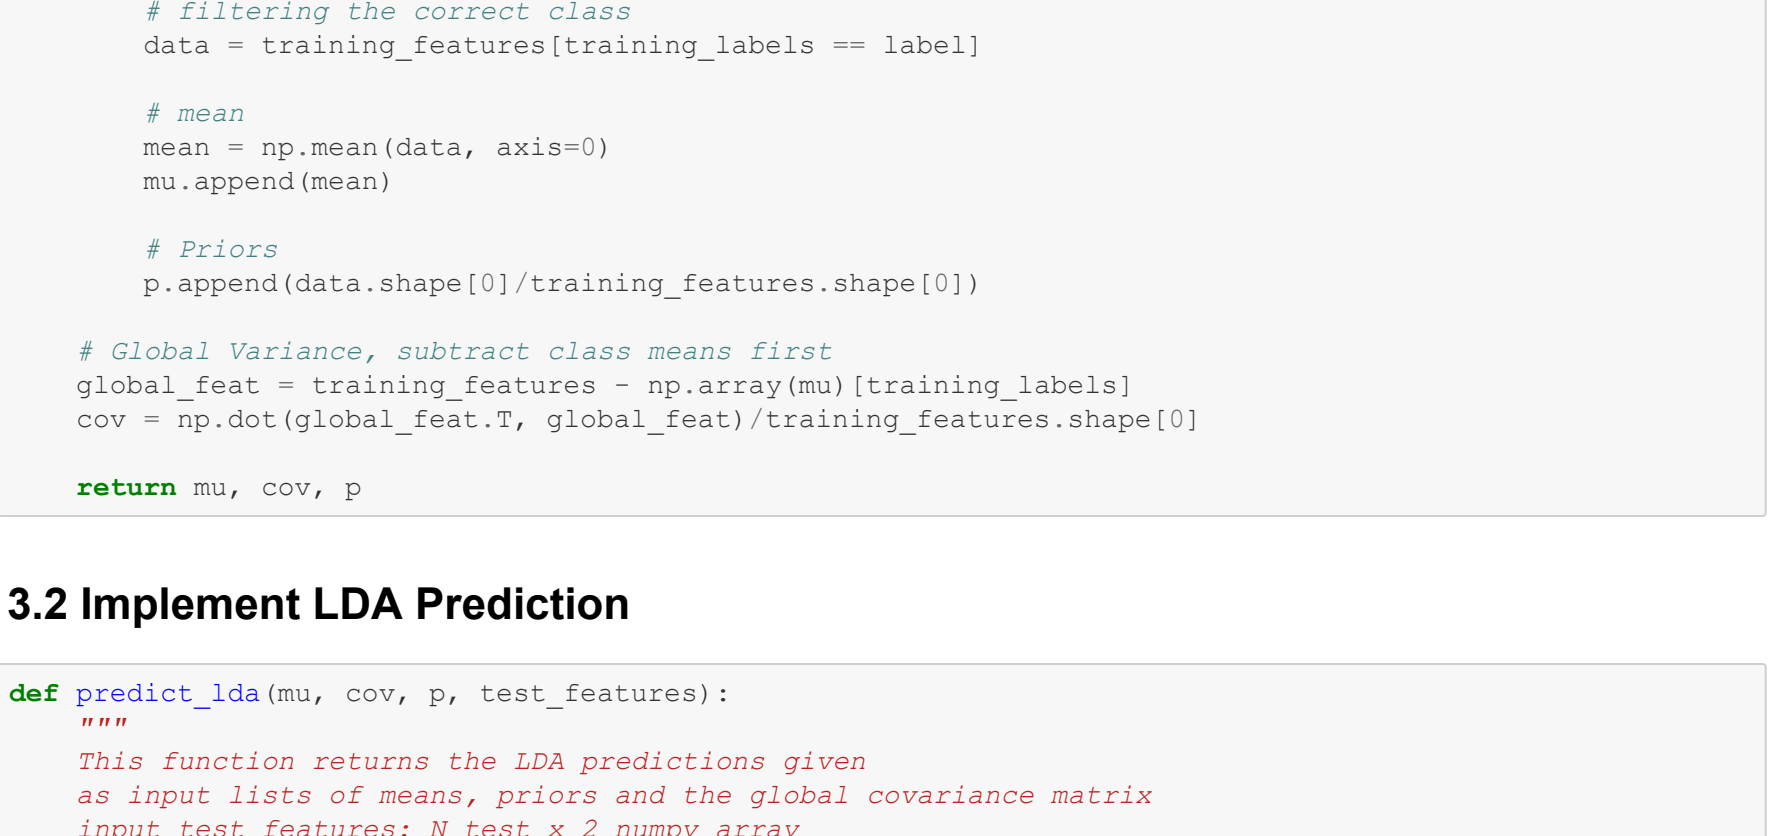
```
In [9]: # Dimension Reduction
_Xr_train, _Xr_test = worse_features2d(X_train), worse_features2d(X_test)

# Scatter Plot
scatter_plot_simple(_Xr_train, y_train, "Training - Worse features")
```



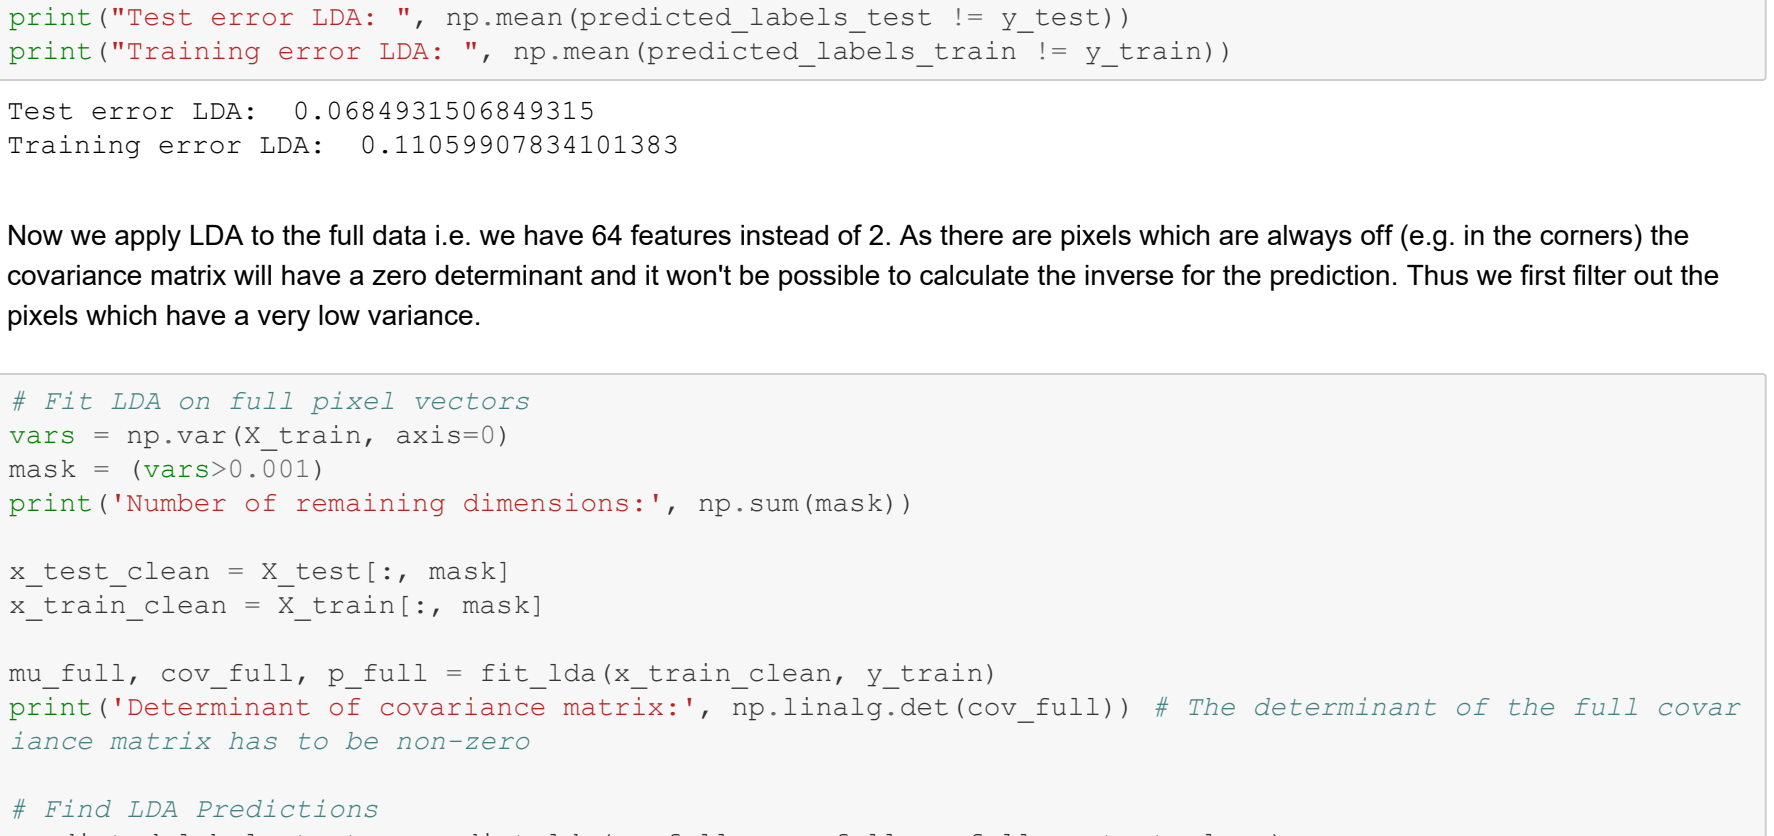
```
In [9]: # Dimension Reduction
_Xr_train, _Xr_test = worse_features2d(X_train), worse_features2d(X_test)

# Scatter Plot
scatter_plot_simple(_Xr_train, y_train, "Training - Worse features")
```



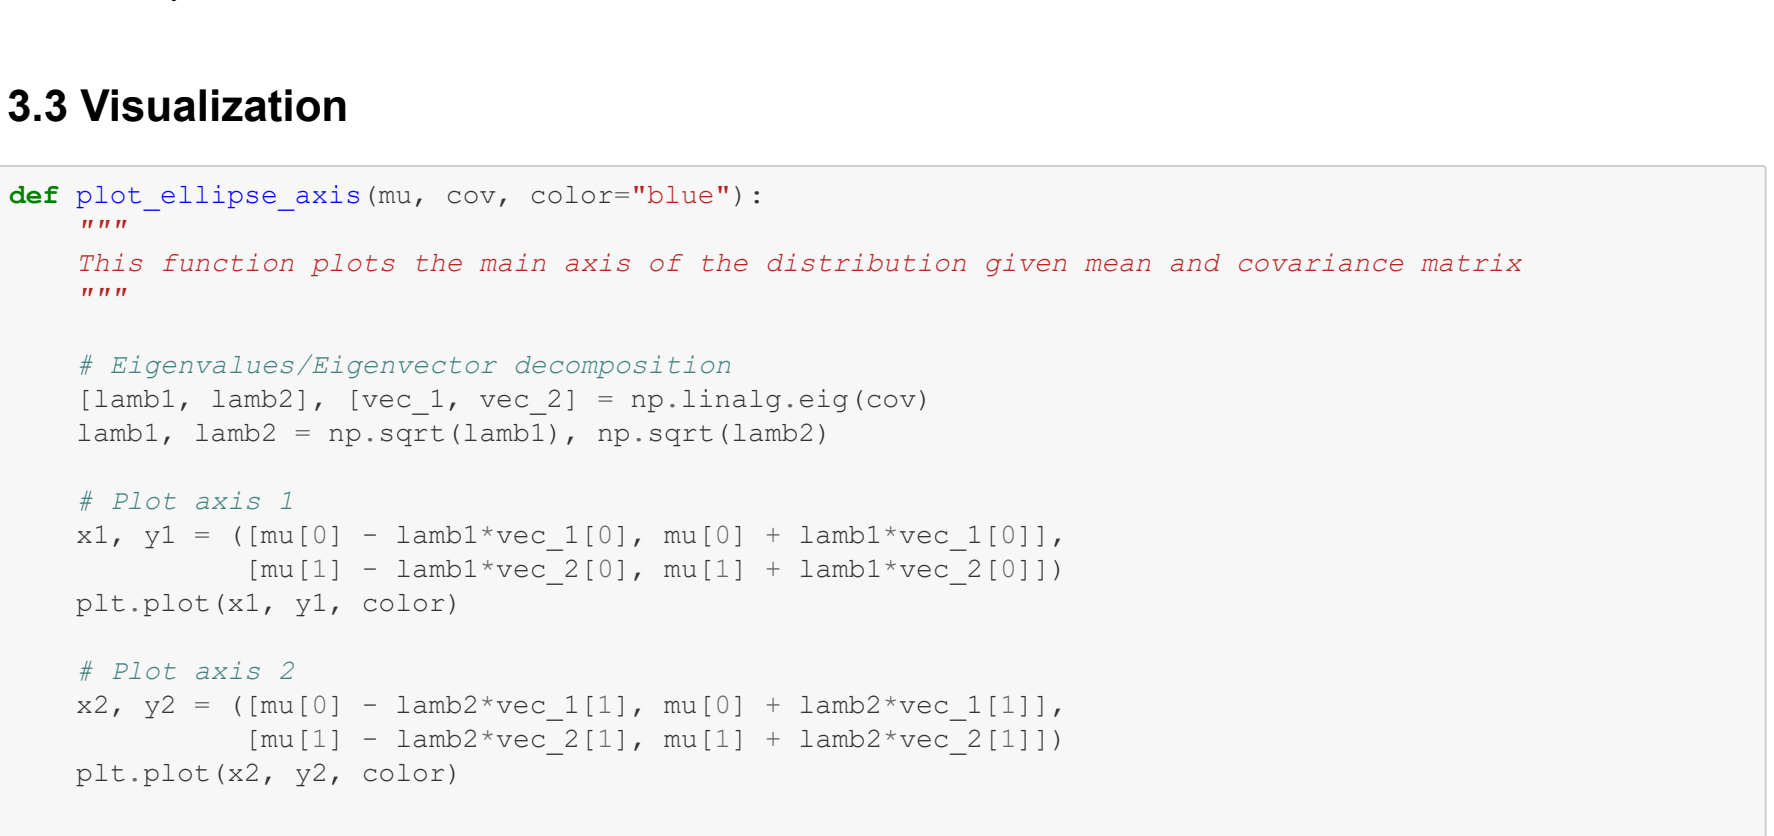
```
In [9]: # Dimension Reduction
_Xr_train, _Xr_test = worse_features2d(X_train), worse_features2d(X_test)

# Scatter Plot
scatter_plot_simple(_Xr_train, y_train, "Training - Worse features")
```



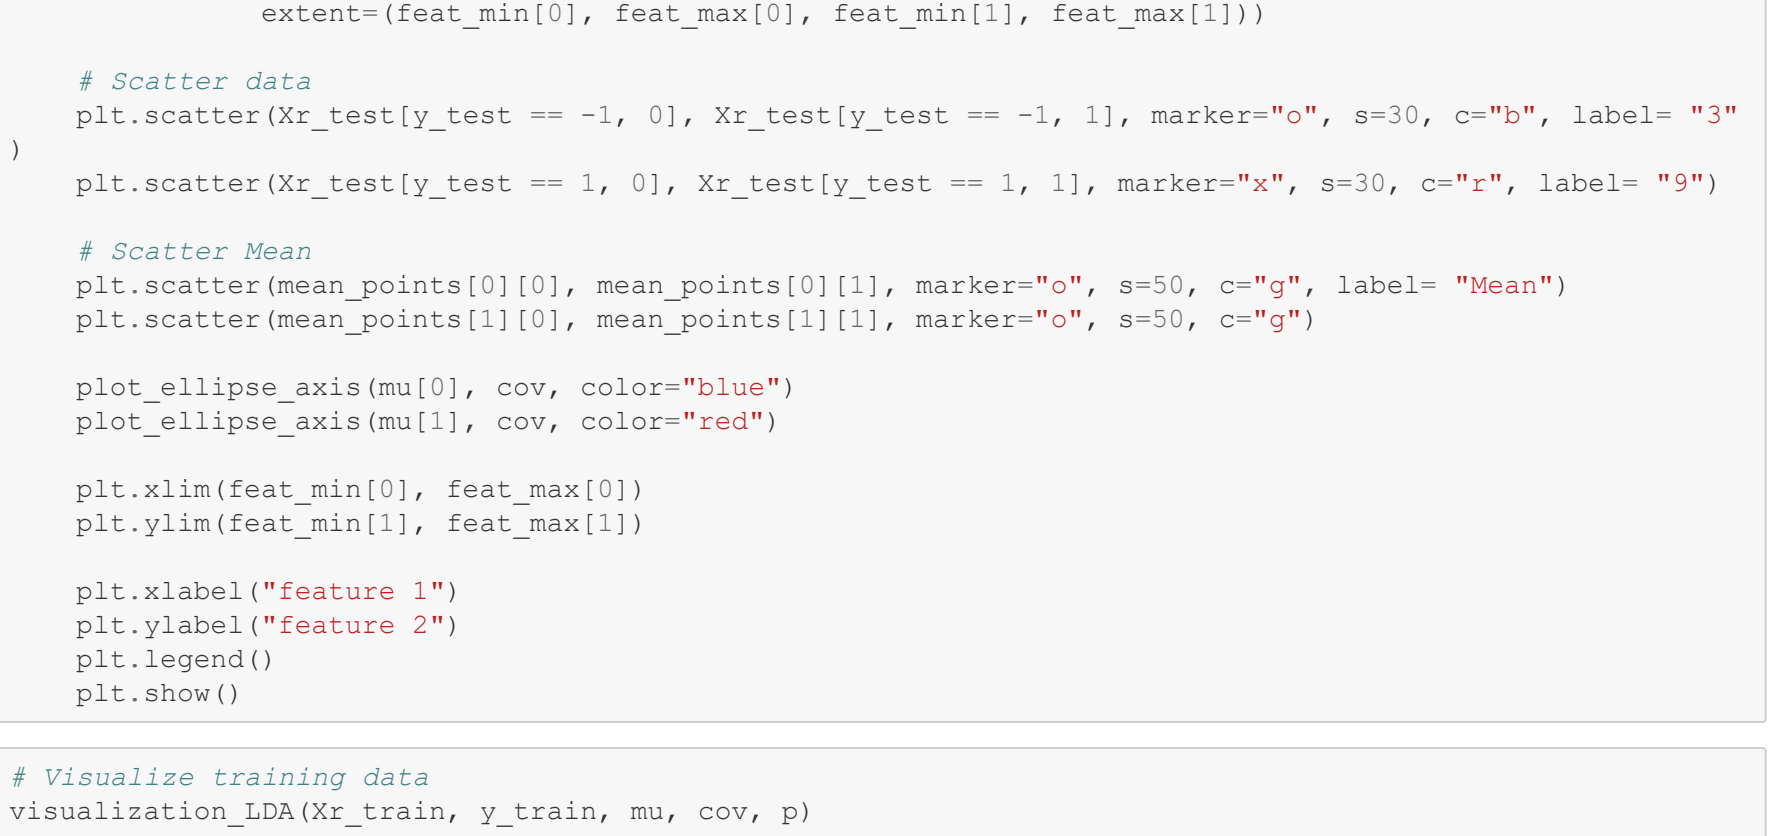
```
In [9]: # Dimension Reduction
_Xr_train, _Xr_test = worse_features2d(X_train), worse_features2d(X_test)

# Scatter Plot
scatter_plot_simple(_Xr_train, y_train, "Training - Worse features")
```



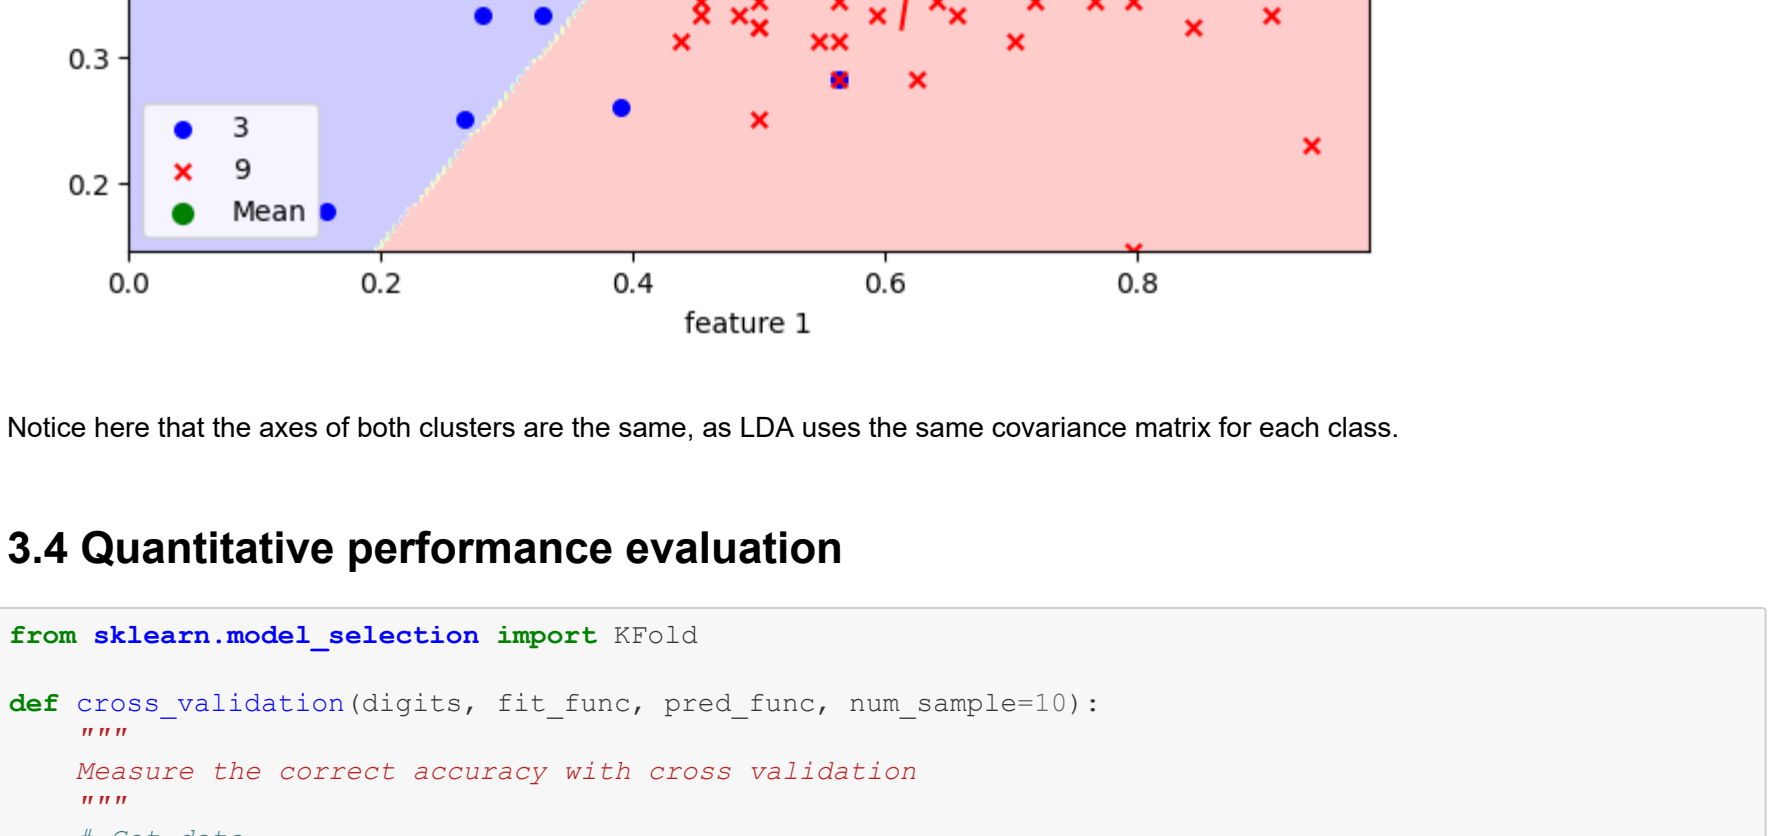
```
In [9]: # Dimension Reduction
_Xr_train, _Xr_test = worse_features2d(X_train), worse_features2d(X_test)

# Scatter Plot
scatter_plot_simple(_Xr_train, y_train, "Training - Worse features")
```



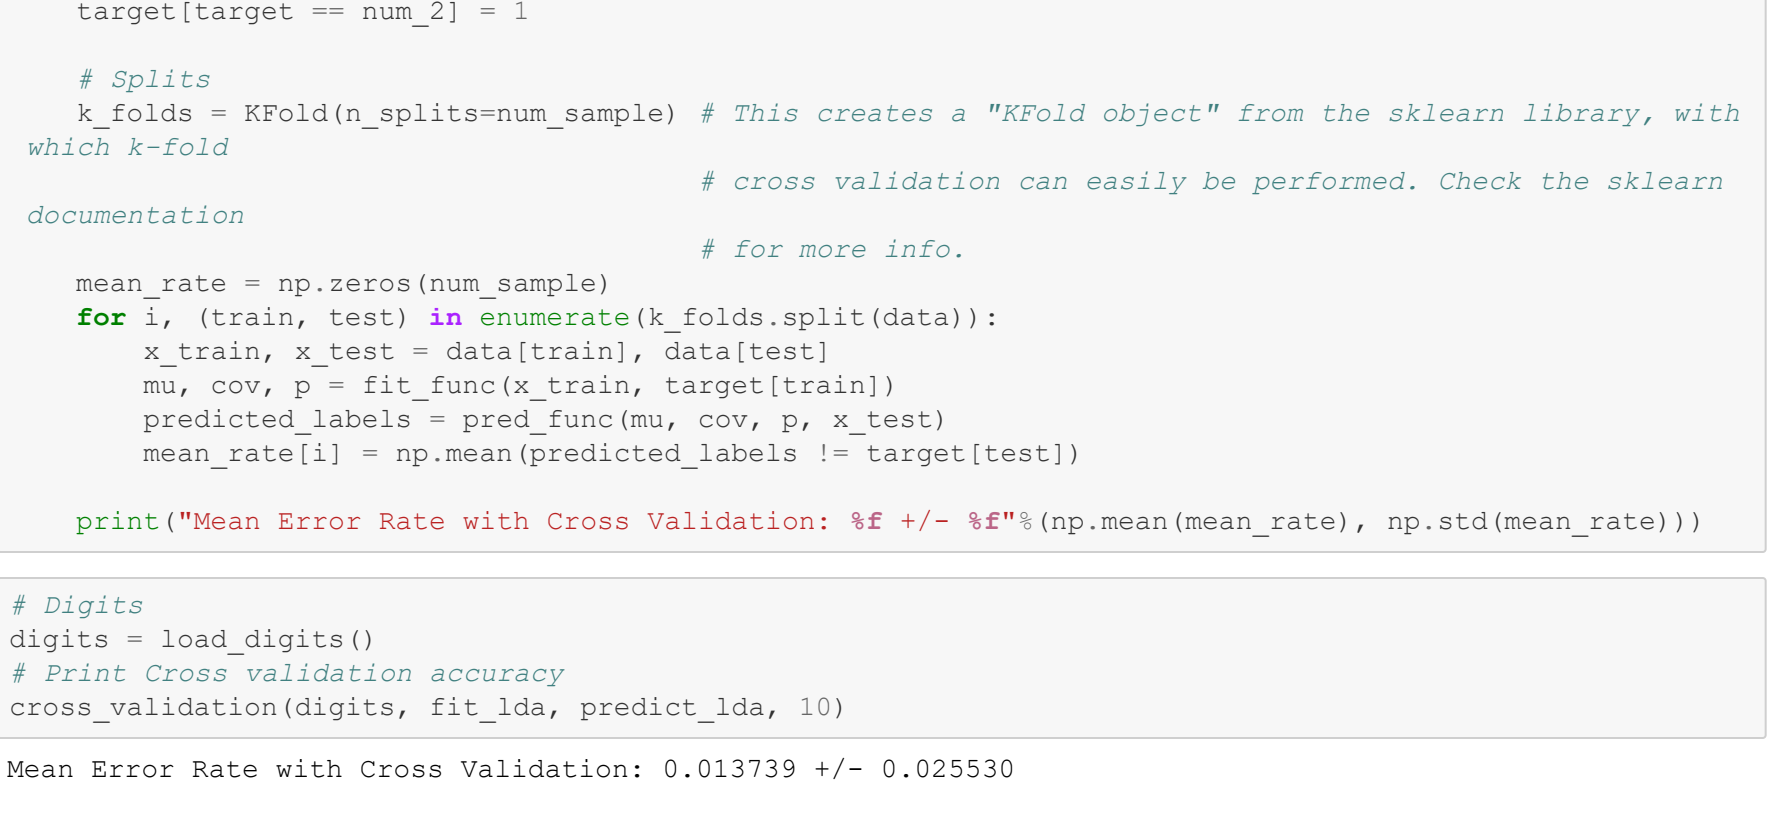
```
In [9]: # Dimension Reduction
_Xr_train, _Xr_test = worse_features2d(X_train), worse_features2d(X_test)

# Scatter Plot
scatter_plot_simple(_Xr_train, y_train, "Training - Worse features")
```




```
In [9]: # Dimension Reduction
_Xr_train, _Xr_test = worse_features2d(X_train), worse_features2d(X_test)

# Scatter Plot
scatter_plot_simple(_Xr_train, y_train, "Training - Worse features")
```



```
In [9]: # Dimension Reduction
_Xr_train, _Xr_test = worse_features2d(X_train), worse_features2d(X_test)

# Scatter Plot
scatter_plot_simple(_Xr_train, y_train, "Training - Worse features")
```

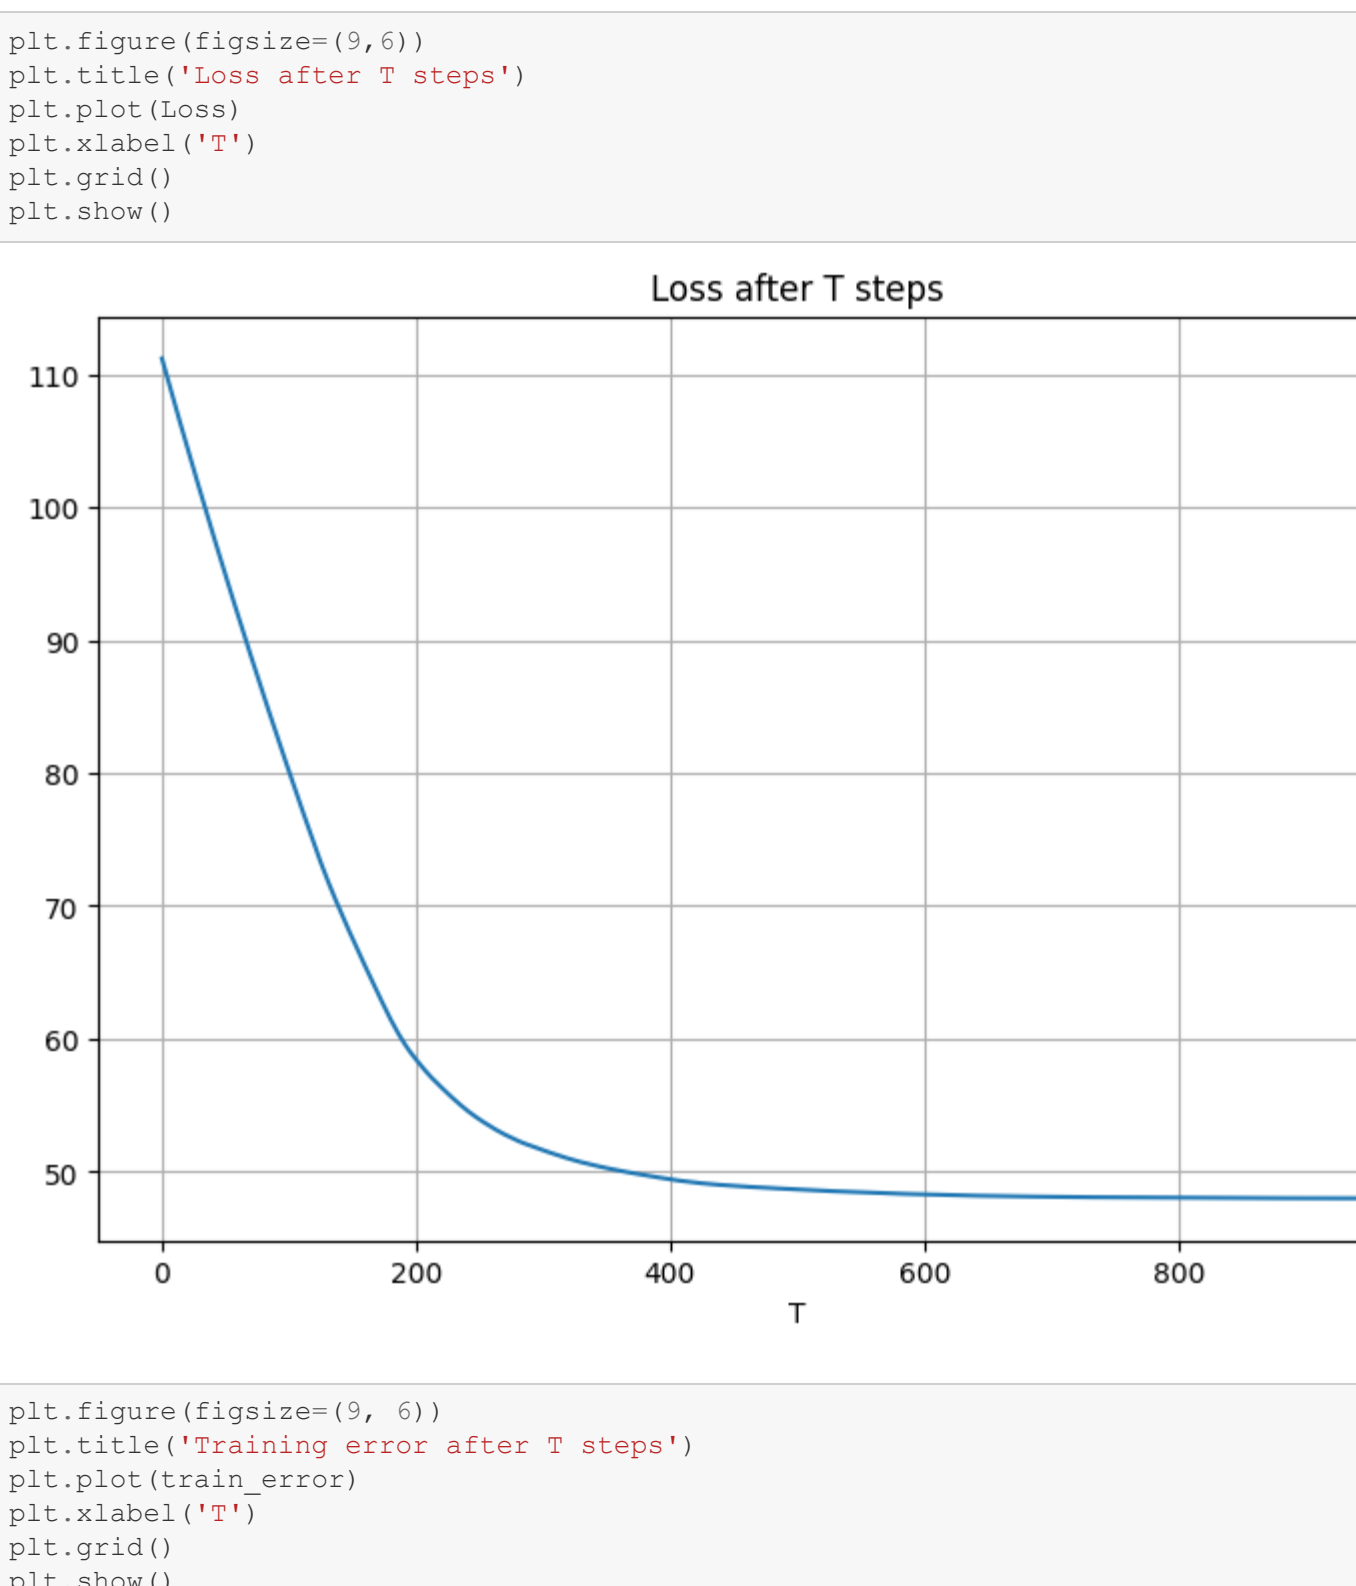




```
In [34]: def fit_SVM(features, training_labels, l=1, T=1000, tau=0.01):  
    """  
    input training_features: N_training x N_features numpy array  
    input training_labels: N_training x 1 numpy array  
    input l: float, T: int, Tau: float  
    output beta, b, list of Loss and train_error  
    """  
  
    # Initial guess for beta and b  
    beta = np.random.normal(loc=0, scale=1, size=(training_features.shape[1]))  
    b = 0  
  
    def ReLU(x):  
        return np.where(x>0, x, 0)  
  
    Loss = []  
    train_error = []  
  
    N = len(training_labels)  
  
    # Implement training loop using the gradient descent algorithm  
    for i in range(T):  
        # Helper variable  
        temp = training_labels*(training_features @ beta + b)  
        # Compute Loss  
        Loss.append(1/2*np.dot(beta, beta) + 1*np.mean(ReLU(1 - temp)))  
        # Compute training error at each iteration with current guess of beta and b  
        predicted_labels_train = predict_SVM(beta, b, training_features)  
        train_error.append(np.mean(predicted_labels_train != training_labels))  
        # Compute gradients of beta and b  
        grad_beta = beta + 1*np.mean(np.where(temp[1,:None] < 1, -training_labels[1, :None])*training_features,  
                                0), axis=0)  
        grad_b = 1*np.mean(np.where(temp < 1, -training_labels, 0))  
        # Update current guess of beta and b  
        beta = beta - tau*grad_beta  
        b = b - tau*grad_b  
  
    return beta, b, Loss, train_error
```

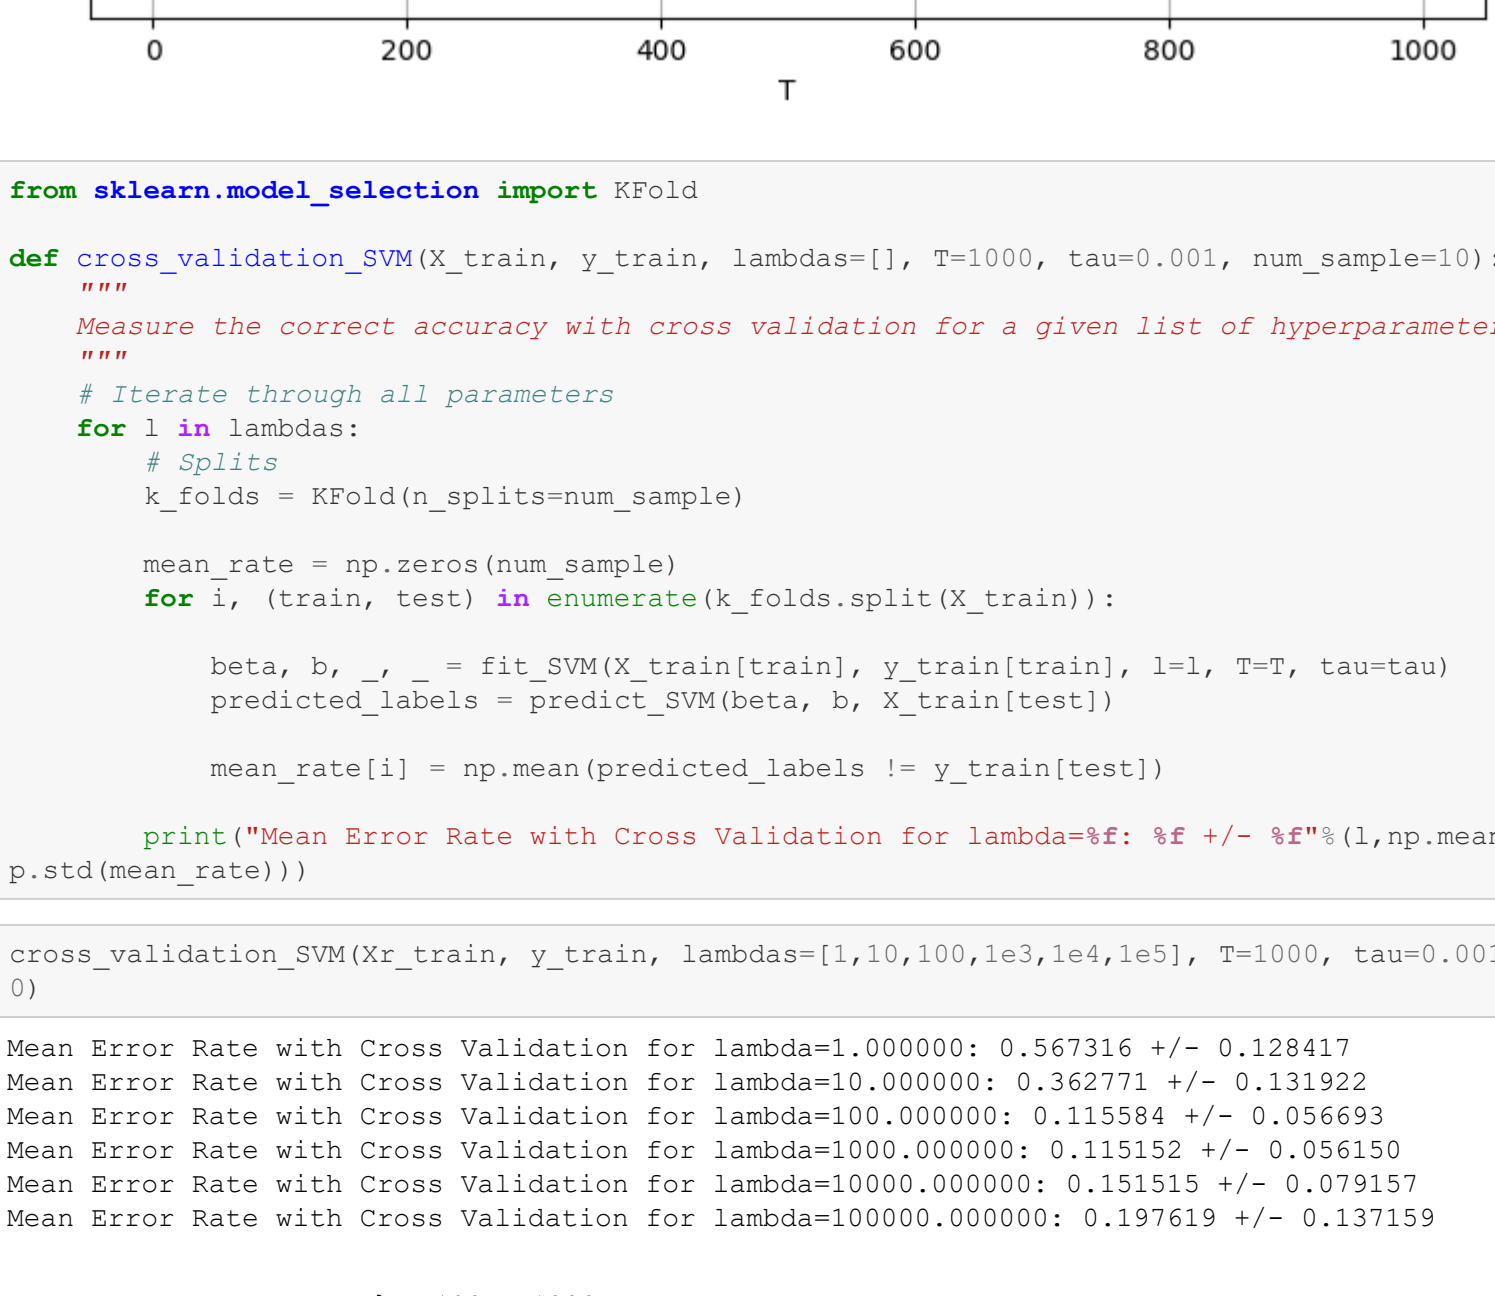
```
In [43]: # Fit SVM with hand-picked hyperparameters  
beta, b, Loss, train_error = fit_SVM(Xr_train, y_train, l=100, T=1000, tau=0.001)  
# Print beta and b  
print("beta =", beta)  
print("b =", b)  
  
beta = [ 4.20689925 -1.58718237]  
b = -1.0986175115207266
```

```
In [44]: # Visualizations for SVM  
visualization_SimpleLinearClassifier(beta, b, Xr_test, y_test)
```

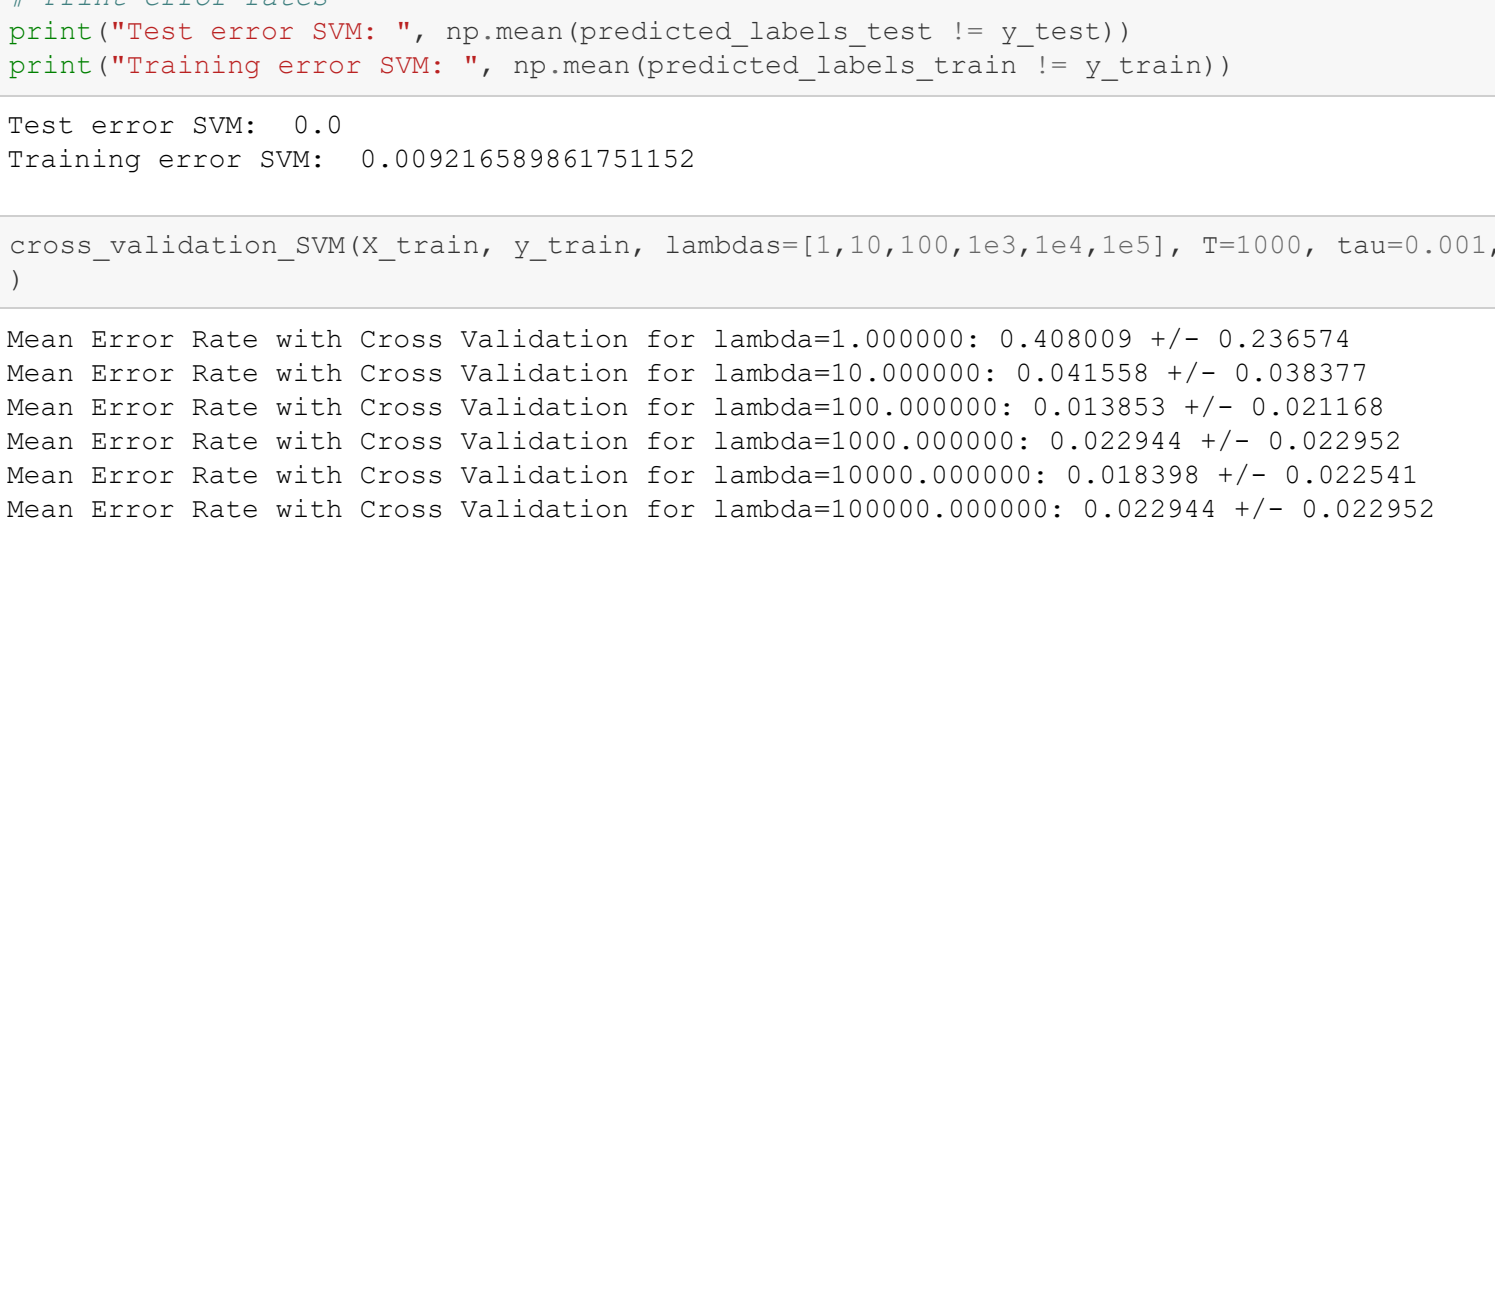


```
In [45]: # Find SVM Predictions  
predicted_labels_test = predict_SVM(beta, b, Xr_test)  
predicted_labels_train = predict_SVM(beta, b, Xr_train)  
  
# Print error rates  
print("Test error SVM: ", np.mean(predicted_labels_test != y_test))  
print("Training error SVM: ", np.mean(predicted_labels_train != y_train))  
  
Test error SVM: 0.0410938904105658  
Training error SVM: 0.10138248847926268
```

```
In [53]: plt.figure(figsize=(9,6))  
plt.title('Loss after T steps')  
plt.plot(Loss)  
plt.xlabel('T')  
plt.grid()  
plt.show()
```



```
In [54]: plt.figure(figsize=(9, 6))  
plt.title('Training error after T steps')  
plt.plot(train_error)  
plt.xlabel('T')  
plt.grid()  
plt.show()
```



```
In [55]: from sklearn.model_selection import KFold  
  
def cross_validation_SVM(X_train, y_train, lambdas=[], T=1000, tau=0.001, num_sample=10):  
    """  
    Measure the correct accuracy with cross validation for a given list of hyperparameters (lambdas)  
    """  
    # Iterate through all parameters  
    for l in lambdas:  
        # Splits  
        K_folds = KFold(n_splits=num_sample)  
  
        mean_rate = np.zeros(num_sample)  
  
        for i, (train, test) in enumerate(K_folds.split(X_train)):  
            beta, b, _ = fit_SVM(X_train[train], y_train[train], l=l, T=T, tau=tau)  
            predicted_labels = predict_SVM(beta, b, X_train[test])  
  
            mean_rate[i] = np.mean(predicted_labels != y_train[test])  
  
        print("Mean Error Rate with Cross Validation for lambda=%f: %f +/- %f"%(l,np.mean(mean_rate), n  
        p.std(mean_rate)))
```

```
In [56]: cross_validation_SVM(Xr_train, y_train, lambdas=[1,10,100,1e3,1e4,1e5], T=1000, tau=0.001, num_sample=1  
0)  
  
Mean Error Rate with Cross Validation for lambda=1.000000: 0.567316 +/- 0.126417  
Mean Error Rate with Cross Validation for lambda=10.000000: 0.362771 +/- 0.131922  
Mean Error Rate with Cross Validation for lambda=100.000000: 0.115584 +/- 0.056693  
Mean Error Rate with Cross Validation for lambda=1000.000000: 0.115152 +/- 0.056150  
Mean Error Rate with Cross Validation for lambda=10000.000000: 0.151515 +/- 0.079157  
Mean Error Rate with Cross Validation for lambda=100000.000000: 0.197619 +/- 0.137159  
  
Best results are achieved with  $\lambda = 100$  or  $1000$ .
```

```
And now once again with the full pixel vector.
```

```
In [57]: # Fit SVM on full pixel data  
beta, b, Loss, train_error = fit_SVM(X_train, y_train, l=100, T=1000, tau=0.001)  
  
# Find SVM Predictions  
predicted_labels_test = predict_SVM(beta, b, X_test)  
predicted_labels_train = predict_SVM(beta, b, X_train)  
  
# Print error rates  
print("Test error SVM: ", np.mean(predicted_labels_test != y_test))  
print("Training error SVM: ", np.mean(predicted_labels_train != y_train))  
  
Test error SVM: 0.0  
Training error SVM: 0.009216589861751152
```

```
In [58]: cross_validation_SVM(X_train, y_train, lambdas=[1,10,100,1e3,1e4,1e5], T=1000, tau=0.001, num_sample=10  
)  
  
Mean Error Rate with Cross Validation for lambda=1.000000: 0.488089 +/- 0.236574  
Mean Error Rate with Cross Validation for lambda=10.000000: 0.041558 +/- 0.058377  
Mean Error Rate with Cross Validation for lambda=100.000000: 0.013853 +/- 0.021168  
Mean Error Rate with Cross Validation for lambda=1000.000000: 0.022944 +/- 0.022952  
Mean Error Rate with Cross Validation for lambda=10000.000000: 0.018398 +/- 0.022561  
Mean Error Rate with Cross Validation for lambda=100000.000000: 0.022944 +/- 0.022952
```